

COMP4220: Machine Learning, Spring 2022, Assignment 5

Please submit one pdf file for all questions.

1. List five hyperparameters you can tweak in a basic neural network?

Numbers of layers, number of neurons per layer, activation functions, learning rate, and batch size.

2. What is backpropagation and how does it work?

Going from the output layer to the input layer of neural network to optimize weights to minimize loss function using the chain rule.

Programming Assignment (Artificial Neural Network-ANN)

Importing libraries

```
import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.compose import ColumnTransformer
import keras
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, precision_score,
accuracy_score, f1_score, recall_score
import matplotlib.pyplot as plt
```

*# Importing the dataset. This dataset describes churning, which is
the rate at which customers stop doing business with a company*
dataset = pd.read_csv('Churn_Modelling.csv')
dataset

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
0	1	15634602	Hargrave	619	France	Female
42						
1	2	15647311	Hill	608	Spain	Female
41						
2	3	15619304	Onio	502	France	Female
42						
3	4	15701354	Boni	699	France	Female

```

39
4          5      15737888      Mitchell      850      Spain      Female
43
...      ...      ...      ...      ...      ...      ...
...
9995      9996      15606229      Obijiaku      771      France      Male
39
9996      9997      15569892      Johnstone      516      France      Male
35
9997      9998      15584532      Liu      709      France      Female
36
9998      9999      15682355      Sabbatini      772      Germany      Male
42
9999      10000      15628319      Walker      792      France      Female
28

```

```

      Tenure      Balance      NumOfProducts      HasCrCard      IsActiveMember      \
0          2          0.00          1          1          1
1          1      83807.86          1          0          1
2          8     159660.80          3          1          0
3          1          0.00          2          0          0
4          2     125510.82          1          1          1
...      ...      ...      ...      ...      ...
9995      5          0.00          2          1          0
9996     10      57369.61          1          1          1
9997      7          0.00          1          0          1
9998      3      75075.31          2          1          0
9999      4     130142.79          1          1          0

```

```

      EstimatedSalary      Exited
0          101348.88          1
1          112542.58          0
2          113931.57          1
3          93826.63          0
4          79084.10          0
...      ...      ...
9995          96270.64          0
9996         101699.77          0
9997          42085.58          1
9998          92888.52          1
9999          38190.78          0

```

[10000 rows x 14 columns]

1. Looking at the dataset we can see that the first 3 columns are not essential for our model.

Make a X variable that contains all other columns except the first three columns and Exited (label) Make a Y variable (the Exited column)

```
X = dataset.iloc[:, 3:13].values # Select input features X
y = dataset.iloc[:, 13].values   # The last column "Exited" is the
output variable Y
```

```
print(X.shape, y.shape)
```

```
(10000, 10) (10000,)
```

2. In X there are Geography and Gender columns that are in string format which we can't use for training. Thus we should transform them into numerical type to train our model.

Use LabelEncoder and OneHotEncoder from sklearn.preprocessing to transform the "Geography" and "Gender" columns into numerical data type

```
# Encoding categorical data
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
labelencoder_X_1 = LabelEncoder()
X[:, 1] = labelencoder_X_1.fit_transform(X[:, 1])
labelencoder_X_2 = LabelEncoder()
X[:, 2] = labelencoder_X_2.fit_transform(X[:, 2])
ct = ColumnTransformer([("Geography", OneHotEncoder(), [1])],
remainder = 'passthrough')
X = ct.fit_transform(X)
```

3. Split the dataset into the Training set and Test set (test_size = 0.2)

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, stratify=y, random_state=0)
```

4. Apply Feature Scaling to all features before training a neural network

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
X_train = pd.DataFrame(scaler.fit_transform(X_train))
X_test = pd.DataFrame(scaler.transform(X_test))
```

5. Let's build ANN model by using the Keras sequential package

Initialize the sequential model Add the input layer and the first hidden layer Hint:
For the first layer use (units = 6, kernel_initializer = 'uniform', activation = 'relu',
input_dim = 11)

```
import keras
from keras.models import Sequential
from keras.layers import Dense

classifier = Sequential()
classifier.add(Dense(units = 6, kernel_initializer = 'uniform',
activation = 'relu'))
```

6. Add the second hidden layer

Hint: (units = 6, kernel_initializer = 'uniform', activation = 'relu')

```
classifier.add(Dense(units = 6, kernel_initializer = 'uniform',  
activation = 'relu'))
```

7. Add the output layer

Hint: (units = 1, kernel_initializer = 'uniform', activation = 'sigmoid')

```
classifier.add(Dense(units = 1, kernel_initializer = 'uniform',  
activation = 'sigmoid'))
```

8. Compile the ANN

hint: (optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])

```
classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy',  
metrics = ['accuracy'])
```

9. Fit the ANN to the training set

(batch_size = 5, epochs = 20)

```
batch_size = 5
```

```
epochs = 20
```

```
#history = classifier.fit()
```

```
classifier.fit(X_train, y_train, batch_size = 5, epochs = 20)
```

Epoch 1/20

1600/1600 [=====] - 3s 2ms/step - loss: 0.4699 - accuracy: 0.7956

Epoch 2/20

1600/1600 [=====] - 3s 2ms/step - loss: 0.4163 - accuracy: 0.7962

Epoch 3/20

1600/1600 [=====] - 3s 2ms/step - loss: 0.4035 - accuracy: 0.8135

Epoch 4/20

1600/1600 [=====] - 3s 2ms/step - loss: 0.3920 - accuracy: 0.8303

Epoch 5/20

1600/1600 [=====] - 3s 2ms/step - loss: 0.3828 - accuracy: 0.8371

Epoch 6/20

1600/1600 [=====] - 3s 2ms/step - loss: 0.3762 - accuracy: 0.8421

Epoch 7/20

1600/1600 [=====] - 3s 2ms/step - loss: 0.3707 - accuracy: 0.8466

Epoch 8/20

1600/1600 [=====] - 3s 2ms/step - loss: 0.3669 - accuracy: 0.8519

```

Epoch 9/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3629 - accuracy: 0.8508
Epoch 10/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3604 - accuracy: 0.8529
Epoch 11/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3584 - accuracy: 0.8534
Epoch 12/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3576 - accuracy: 0.8561
Epoch 13/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3557 - accuracy: 0.8571
Epoch 14/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3555 - accuracy: 0.8561
Epoch 15/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3532 - accuracy: 0.8585
Epoch 16/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3536 - accuracy: 0.8585
Epoch 17/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3521 - accuracy: 0.8583
Epoch 18/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3525 - accuracy: 0.8571
Epoch 19/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3513 - accuracy: 0.8575
Epoch 20/20
1600/1600 [=====] - 3s 2ms/step - loss:
0.3505 - accuracy: 0.8584

```

<keras.callbacks.History at 0x7fc566f64090>

10. Make predictions and evaluate the model

hint: just consider `y_pred` the values where `y_pred` is greater than 0.5 (`y_pred = (y_pred > 0.5)`) Make the confusion matrix and show the result Evaluate the precision, accuracy, recall, and f1 score and show the result

```

y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5)

```

```

from sklearn.metrics import confusion_matrix, accuracy_score,
f1_score, precision_score, recall_score
acc = accuracy_score(y_test, y_pred)

```

```
matrix = pd.DataFrame([[ 'ANN', acc]],  
                       columns = [ 'Model', 'Accuracy'])
```

```
matrix
```

```
   Model  Accuracy  
0    ANN    0.8545
```

11. Compute the accuracy, precision, recall, and f1 score

```
print('Accuracy: {}'.format(accuracy_score(y_test, y_pred)))  
print('Precision: {}'.format(precision_score(y_test, y_pred)))  
print('Recall: {}'.format(recall_score(y_test, y_pred)))  
print('F1 Score: {}'.format(f1_score(y_test, y_pred)))
```

```
Accuracy: 0.8545  
Precision: 0.8186813186813187  
Recall: 0.36609336609336607  
F1 Score: 0.5059422750424448
```

12. Using Tensorflow Playground

Visit the TensorFlow Playground at <https://playground.tensorflow.org/>

Spend some time playing with this UI to grow your intuition about neural networks. Complete the following problems in a single sitting please.

1. Layers and patterns: try training the default neural network by clicking the run button (top left). Notice how it quickly finds a good solution for the classification task. Notice that the neurons in the first hidden layer have learned simple patterns, while the neurons in the second hidden layer have learned to combine the simple patterns of the first hidden layer into more complex patterns. What happens when you add more layers?

Adding more layers increases the number of weights in the neural network and increases the complexity

1. Activation function: try replacing the Tanh activation function with the ReLU activation function, and train the network again. Notice that it finds a solution even faster, but this time the boundaries are linear. What about the ReLU function causes this?

The boundaries are linear because it is linear for $X > 0$, but non-differentiable at $Z = 0$.

1. Local minima: modify the network architecture to have just one hidden layer with three neurons. Train it multiple times (to reset the network weights, click the reset button next to the play button). What do you notice about the training time?

The training time increases.

1. Too small: now remove one neuron to keep just 2. Notice that the neural network is now incapable of finding a good solution, even if you try multiple times. What do you observe about the number of parameters and the training set?

The number of parameters decreases since the number of neurons per layer is a hyperparameter and causes underfitting.

1. Large enough: next, set the number of neurons to 8 and train the network several times. Notice that it is now consistently fast and never gets stuck. What do you observe about local minima?

Less often of a local minima which means less chances of it getting stuck